

# OeNB Industry Lab - Documentation

Tristan Leiter

November 7, 2025

# Contents

<b>1</b>	<b>Einleitung</b>	<b>2</b>
1.1	Einleitung . . . . .	2
<b>2</b>	<b>Exploratory Data Analysis</b>	<b>3</b>
2.1	Overview . . . . .	3
2.1.1	Description of the dataset . . . . .	3
2.1.2	Description of the dataset . . . . .	3
2.2	Data Splitting with Imbalanced Data . . . . .	3
2.2.1	Class Imbalance . . . . .	3
2.2.2	The Solution: Stratified Sampling . . . . .	3
<b>3</b>	<b>Loss-function</b>	<b>4</b>
3.1	Area under the curve (AuC) . . . . .	4
3.2	Recall . . . . .	4
<b>4</b>	<b>Generalized linear models (GLM)</b>	<b>5</b>
4.1	Overview . . . . .	5
<b>5</b>	<b>Regularized GLMs</b>	<b>6</b>
5.1	Description of the algorithm . . . . .	6
5.1.1	Description of the algorithm . . . . .	7
<b>A</b>	<b>Overview</b>	<b>8</b>

# Chapter 1

## Einleitung

### 1.1 Einleitung

# Chapter 2

## Exploratory Data Analysis

### 2.1 Overview

#### 2.1.1 Description of the dataset

#### 2.1.2 Description of the dataset

### 2.2 Data Splitting with Imbalanced Data

#### 2.2.1 Class Imbalance

In this credit risk dataset, the target variable,  $y$ , is inherently imbalanced. The number of non-defaults (0) significantly outnumbers the number of defaults (1). This is a common and expected characteristic of credit risk data.

This imbalance poses a significant challenge for model development. If we were to use a simple random split to create our training, validation, and test sets, we would face a high risk of creating unrepresentative samples. For example, a small test set could, purely by chance, end up with a much higher or lower percentage of defaults than the original dataset—or, in the worst case, zero defaults.

#### 2.2.2 The Solution: Stratified Sampling

To prevent this, we employ stratified sampling. This is a technique that ensures the original class distribution of the target variable is preserved in each of the new data splits.

Here is the reasoning for its use:

**Guarantees Representation:** Stratification forces the splits to maintain the original ratio of defaults to non-defaults. If 0.0865% of the original dataset are defaults, the training set, validation set, and test set will all contain approximately 0.0865% defaults.

**Enables Reliable Evaluation:** When the test set is representative, the performance metrics we calculate (like accuracy, precision, recall, and F1-score) are meaningful. Evaluating a model on a test set with a skewed default rate would give us a misleading and over-optimistic (or pessimistic) score.

**Promotes Model Generalization:** By training the model on a set that accurately reflects the real-world data distribution, we help it learn the patterns of both the majority (non-default) and minority (default) classes, leading to a more robust and generalizable model.

In summary, using stratified sampling on the  $y$  variable is a critical step to ensure our model is trained and evaluated on a reliable, representative foundation.

# Chapter 3

## Loss-function

### 3.1 Area under the curve (AuC)

The \*\*Area Under the Curve (AUC)\*\* is a single, aggregate metric that evaluates the performance of a binary classification model across all possible classification thresholds.

It is the area under the \*\*ROC (Receiver Operating Characteristic) curve\*\* , which plots the model's **Sensitivity** (True Positive Rate) against its **Specificity** (True Negative Rate) at every conceivable threshold.

- An **AUC of 1.0** represents a perfect model that can distinguish between positive and negative classes with 100% accuracy.
- An **AUC of 0.5** represents a model with no discriminatory power, equivalent to a random guess (as shown by the diagonal line in the plot).
- A good model will have an AUC well above 0.5.

**Relevance to Credit Risk:** The AUC is extremely valuable because it measures the model's ability to *rank* clients by risk. It answers the question: "If I pick a random defaulting firm and a random non-defaulting firm, what is the probability that my model gives a higher risk score to the defaulting firm?"

A high AUC (e.g., > 0.75) means the model is effective at separating the "bad" clients from the "good" ones. This is crucial before deciding on a specific *business-level threshold* (like "what probability score triggers a loan rejection?"). It tells us the model itself is fundamentally sound.

### 3.2 Recall

Recall, also known as **Sensitivity** or the **True Positive Rate (TPR)**, is a metric that measures a model's ability to identify all relevant instances of a class.

It is calculated using the following formula, based on the outputs of a **confusion matrix** at a specific threshold:

$$\text{Recall} = \frac{\text{True Positives (TP)}}{\text{True Positives (TP)} + \text{False Negatives (FN)}}$$

- **True Positives (TP):** The number of firms that *actually defaulted* and were *correctly flagged* as defaults by the model.
- **False Negatives (FN):** The number of firms that *actually defaulted* but were *incorrectly flagged* as non-defaults. This is the most costly error in credit risk.

**Relevance to Credit Risk:** Recall answers the single most important question for a risk-averse institution: "**Of all the firms that actually defaulted, what percentage did we successfully catch?**"

In credit risk, the cost of a **False Negative** (missing a default) is extremely high, as it results in a direct financial loss. The cost of a **False Positive** (flagging a healthy firm as risky) is much lower—it might lead to a manual review or a rejected loan application (an opportunity cost).

Therefore, banks often prioritize high Recall. They are willing to accept a higher number of false alarms to ensure they minimize the number of defaults that slip through undetected. Setting a threshold to achieve a specific Recall (e.g., 95%) is a common, risk-averse strategy.

## Chapter 4

# Generalized linear models (GLM)

### 4.1 Overview

# Chapter 5

## Regularized GLMs

### 5.1 Description of the algorithm

The regularized estimator is obtained by solving the following optimization problem:

$$\hat{\beta}^{pen} = \arg \min_{\beta} \left\{ -\frac{1}{N} l(\beta, \phi | \mathbf{y}, \mathbf{X}) + \lambda P(\beta) \right\}$$

This objective function consists of two parts. In our logistic regression model, we predict the probability of default for each firm. We assume that the outcome ( $\mathbf{Y}$ ) is a Bernoulli random variable (a special case of the Binomial distribution). Then we model the probability of this specific firm defaulting given their features. This is reflected in the loss function:

#### 1. The Loss Function: $-\frac{1}{N} l(\beta, \phi | \mathbf{y}, \mathbf{X})$

This term measures how well the model fits the data.

- $l(\beta, \phi | \mathbf{y}, \mathbf{X})$  is the \*\*log-likelihood function\*\*. Since we assume a Bernoulli/Binomial outcome, this is the log-likelihood of the binomial distribution.
- We want to find parameters ( $\beta$ ) that *maximize* the likelihood of observing our data.
- Maximizing the log-likelihood is equivalent to *minimizing* the \*\*negative log-likelihood\*\* (which is why the negative sign is there).
- The  $\frac{1}{N}$  term scales it by the number of observations  $N$  to get the \*\*average negative log-likelihood\*\*. This is also known as the **log-loss** or **cross-entropy loss**.

#### 2. The Penalty Term: $\lambda P(\beta)$

This term, where  $\lambda \geq 0$ , penalizes the size of the coefficients to prevent overfitting. The **alpha** parameter controls the type of penalty:

- **alpha = 1** corresponds to the \*\*Lasso\*\* estimator (L1 penalty).
- **alpha = 0** corresponds to the \*\*Ridge\*\* estimator (L2 penalty).

The optimal regularization parameter,  $\lambda$ , is found by  $k$ -fold cross-validation. For example, for  $k = 10$ , the algorithm splits the training data into 10 equal-sized folds. It then iterates 10 times:

- In each iteration, it trains the model on 9 folds and evaluates it on the 1 held-out fold (e.g., train on folds 2-10, test on fold 1).
- This entire process is repeated for a sequence of different  $\lambda$  values.

Finally, the algorithm calculates the average error (e.g., AUC or deviance) across all 10 folds for each  $\lambda$  value. The  $\lambda$  that produces the lowest average error is chosen as the optimal value.

In the **glmnet** package, this is implemented by:

```
cv_model_logit <- cv.glmnet(x_train,
                               y_train,
                               family = "binomial",
                               alpha = 1)
```

The optimal  $\lambda$  value is stored in `cv_model_logit$lambda.min`, which is then used by the `predict()` function.

### 5.1.1 Description of the algorithm

## **Appendix A**

## **Overview**